

Breaking substitution ciphers with Markov models

Algorithms description and implementation notes

Enrico Rollando (e.rollando@outlook.com)

August 6, 2021

Abstract

Cryptography is a useful playground for applying statistical methods. This document uses the simple substitution cipher as a toy problem and shows how to decrypt messages without knowing the key (i.e., without known plaintext). It focuses on Markov model-based approaches, reviews the underlying theory, and provides concrete implementations and usage examples.

Introduction

Can we break a simple substitution cipher automatically? “To break” here means discovering and exploiting weaknesses to restore—at least in part—the original text.

For scope, consider plaintext and ciphertext over uppercase English letters plus space (26+1 symbols). In the base setting, space maps to itself.

A *simple substitution cipher* is a mapping function $S_k(\text{text}) \rightarrow \text{ciphertext}$ where k is the key. It is a bijection on the alphabet (equivalently, a permutation of its symbols). Given the ciphertext only, the goal is to recover the inverse mapping $S_k^{-1}(\text{ciphertext})$ without any known plaintext.

To guide the search, define a function $\text{fit}(\text{text})$ that scores the current guess. Ideally this would answer “how likely is this the original plaintext?”, but a practical proxy is: “how plausible is this as English text?”. This score lets us compare different candidate decryptions and prefer the most plausible one.

A natural way to build such a score is to use language statistics: how often letters and n -grams appear in English, and how likely one character is to follow another. We then score candidate decryptions using these statistics.

There is still a problem: the number of possible keys is huge (26! for the English alphabet), far too many to try exhaustively. Other techniques are needed; this document presents a few of them.

Let us first define the cipher and briefly introduce the language model, which is the main tool used to exploit the cipher’s weakness.

The simple substitution cipher

A simple substitution cipher [1] is a one-to-one mapping between two alphabets: the plaintext alphabet and the ciphertext alphabet. This mapping is the key of the cipher.

Spaces are sometimes removed from plaintext before encryption, as they provide information about word boundaries, which in turn helps recovery (the *Patristocrat* variant).

$P = [a_1, \dots, a_{26}]$, example: ABCDEFGHIJKLMNOPQRSTUVWXYZ
 $C = [c_1, \dots, c_{26}]$, example: HQUVNCFYPXJDIMOBWLAZTGSRKE

$$S_k : P \rightarrow C \tag{1}$$

$$S_k("THISISATEST") = "ZYPAPAHZNAZ"$$

$$S_k^{-1}("ZYPAPAHZNAZ") = "THISISATEST"$$

The Language Model

In any language there are implicit rules we use when we write or interpret sentences. For example, the incomplete sentence “I like to write my letters by ...” is likely to end with “hand”. The preceding words provide context, which makes some continuations more plausible than others.

At the character level, the same idea lets us read an incomplete word like “uncomf.rt.b.y”: plausibility can be interpreted statistically as “which character is most likely here, given its context?”.

A language model is tied to a specific language and captures such regularities using frequencies at the character or word level. As an example, consider a bigram English language model LM_{EN} :

- bigram frequency:
 - how frequent is a bigram ab in English text? $P_{LM}(ab) = \frac{\text{count(occurrences of } ab)}{\text{len(text)} - 1}$
 - for a bigram ab , how frequently does b follow a ?

$$P_{LM}(b | a) = \frac{\text{count}(ab)}{\text{count}(a)} \quad (2)$$

The weakness

A simple substitution cipher always encrypts the same plaintext symbol in the same way. As a result, it preserves a surprising amount of structure from the plaintext: equal letters remain equal, repeated patterns remain repeated, and, more generally, the plaintext n -gram statistics are preserved **up to a relabeling of symbols** induced by the unknown key.

For example, if a ciphertext bigram is a double letter (e.g., XX), then the corresponding plaintext bigram must also be a double letter. By itself this is not enough to identify the letters, but across a whole message these constraints become informative, especially when combined with an English language model.

To exploit this, we can introduce a generative model G for English plaintext, controlled by parameters Θ , and search for a key (or plaintext) that makes the decoded text both (i) consistent with the ciphertext under a substitution mapping and (ii) plausible under G . In practice, this is done iteratively: we refine Θ and/or the key to increase the model score (likelihood or plausibility) of the observed ciphertext under the induced plaintext statistics.

Results at a glance

Method	Cipher length	Iterations/time	Outcome
EM (Baum-Welch)	~120 chars	950 iters (~1-2 min)	Readable plaintext; some letters may need manual swaps
MH (substitution)	~40 chars	50k swaps (~seconds)	Gradual convergence to legible text; prints best candidates
MH (patristocrat)	~30 chars	5k insert/remove (~seconds)	Recovers likely spacing; ambiguity remains on very short texts
Bigram LM prep	200k chars	~2-4 min	Transition matrix A, start dist PI reused across runs

Reproducibility

Goal & Scope

- Ciphertext alphabet: A-Z + space. Spaces are preserved for the base examples; the patristocrat variant removes them.
- Corpus: cleaned English corpus (default: NLTK Brown, `nltk.corpus.brown`)^[7]; preprocessing: uppercase, strip non-letters (except space), collapse whitespace.
- Success criteria: recover readable plaintext and a near-bijective substitution map; report convergence traces instead of only final text.
- Deliverables: pseudocode for EM and Metropolis-Hastings, runnable helpers in `src/crypto`, and reproducible setup (seed + package versions).

How to reproduce

- Inputs: cleaned English corpus (per *Goal & Scope*), ciphertext samples in A-Z + space, patristocrat sample without spaces.
- Runtime (typical laptop): EM bigram LM ~2-4 minutes at 200k chars, EM decryption ~1-2 minutes at 950 iterations; MH 50k swaps on short texts ~seconds; patristocrat spacing ~seconds.
- Seed: fix RNG seeds (e.g., 13/42) for comparable traces; exact ciphertext-dependent results still vary.

Reproducibility setup

- Environment: Python 3.x, `numpy`, optionally `nltk` (only needed if using the Brown corpus).
- Seeding: fix RNG seeds for repeatable EM/MH traces.
- Implementation: place the Python modules from Appendices A-E on your `PYTHONPATH` or in a local `src/crypto` package.
- Corpus preprocessing: uppercase, strip non-letters (except space), collapse whitespace (per *Goal & Scope*)

Methods

Breaking the cipher with Expectation–Maximization (EM / Baum–Welch)

We cast the substitution cipher as a Hidden Markov Model (HMM): hidden plaintext states emit the observed ciphertext symbols, and plaintext-to-plaintext transitions follow an English bigram language model. EM (Baum–Welch) then adjusts the HMM parameters to increase the likelihood of the observed ciphertext. In our setup, we typically learn the transition model (A and π) from an English corpus and, for each ciphertext, re-estimate the emission model (B) while keeping A, π fixed.

Implementation details: scaled forward/backward passes prevent numerical underflow. We also avoid exact zeros in A (e.g., via smoothing) so transitions are never impossible and EM updates do not get stuck.

```
1 # scaled forward/backward to avoid underflow
2 initialize A, B, PI
3 repeat until convergence:
4     alpha, scale = forward(A, B, PI, V)
5     beta = backward(A, B, V, scale)
6     gamma_state = expected_state_visits(alpha, beta)
7     gamma_trans = expected_state_transitions(alpha, beta, A, B, V)
8     if recompute_A: A = normalize_rows(gamma_trans)
9     if recompute_B: B = normalize_rows(state_emissions(gamma_state, V))
10    if recompute_PI: PI = gamma_state[0]
11    B = constrain_space(B) # enforce space->space
```

This highlights the scaled passes, the non-zero constraints on \mathbf{A} , and the bijective tweak for the space character.

EM Implementation: the Baum-Welch algorithm Ciphertext symbols are modeled as emissions from a hidden Markov model (HMM). Let $H = (p_1, \dots, p_N)$ be the hidden state sequence (e.g., plaintext symbols) and $V = (c_1, \dots, c_N)$ the observed sequence (ciphertext symbols). Parameters are $\Theta = (\pi, A, B)$: * $\pi_i = P(p_1 = i)$ initial state probability. * $A_{ij} = P(p_{t+1} = j | p_t = i)$ transition probability. * $B_i(c) = P(c_t = c | p_t = i)$ emission probability.

Joint probability (one particular hidden path):

$$P(V, H | \Theta) = \pi_{p_1} B_{p_1}(c_1) \prod_{t=1}^{N-1} A_{p_t p_{t+1}} B_{p_{t+1}}(c_{t+1})$$

The likelihood $P(V | \Theta) = \sum_H P(V, H | \Theta)$ sums over all possible hidden paths, which is too expensive to do directly. Forward-backward computes it efficiently.

Forward pass (prefix probability). Define:

$$\alpha_t(i) = P(c_{1:t}, p_t = i | \Theta)$$

Initialize and recurse left to right:

$$\begin{aligned} \alpha_1(i) &= \pi_i B_i(c_1) \\ \alpha_t(j) &= B_j(c_t) \sum_i \alpha_{t-1}(i) A_{ij} \quad (t = 2, \dots, N) \end{aligned}$$

Backward pass (suffix probability). Define:

$$\beta_t(i) = P(c_{t+1:N} | p_t = i, \Theta)$$

Start at the end and recurse right-to-left:

$$\begin{aligned} \beta_N(i) &= 1 \\ \beta_t(i) &= \sum_j A_{ij} B_j(c_{t+1}) \beta_{t+1}(j) \quad (t = N-1, \dots, 1) \end{aligned}$$

Likelihood (total probability of the observed ciphertext):

$$P(V | \Theta) = \sum_i \alpha_N(i)$$

Posterior marginals (E-step). Once we have α and β , we can ask: “how likely was the model in state i at time t ?” and “how likely was the transition $i \rightarrow j$ at time t ?”

$$\begin{aligned} \gamma_t(i) &= P(p_t = i | V, \Theta) = \frac{\alpha_t(i) \beta_t(i)}{P(V | \Theta)}, \\ \xi_t(i, j) &= P(p_t = i, p_{t+1} = j | V, \Theta) = \frac{\alpha_t(i) A_{ij} B_j(c_{t+1}) \beta_{t+1}(j)}{P(V | \Theta)} \quad (t = 1, \dots, N-1) \end{aligned}$$

Parameter updates (M-step). Update parameters by normalizing expected counts:

$$\begin{aligned} \pi_i &\leftarrow \gamma_1(i) \\ A_{ij} &\leftarrow \frac{\sum_{t=1}^{N-1} \xi_t(i, j)}{\sum_{t=1}^{N-1} \gamma_t(i)} \end{aligned}$$

$$B_i(c) \leftarrow \frac{\sum_{t=1}^N \mathbf{1}[c_t = c] \gamma_t(i)}{\sum_{t=1}^N \gamma_t(i)}$$

Iterate E/M steps until convergence (or for a fixed number of iterations).

Key implementation details (see [6]): - Transition matrix A avoids exact zeros; strict positivity implies the Markov chain is irreducible and aperiodic. - Forward/backward passes are scaled each step to prevent underflow on long ciphertexts. - First-order LM keeps sparsity manageable; higher orders need smoothing and more data.

EM implementation: a concrete example So far we have summarized the algorithm and its Forward-Backward passes. Let's use it to automatically decrypt some text.

EM usage sketch (see Appendix A for full code)

1. Build bigram LM (A, B_{id}, π) [7]: `A, B_id, PI = build_bigram_model(corpus_text, max_iter=50, limit=200k)`.
2. Initialize emission matrix B randomly, apply `constrain_space_mapping(B)`.
3. Run `BaumWelch.process(A, B, PI, V, max_iter=200, recompute_A=False, recompute_PI=False, ext_fun=constrain_space_mapping)`.
4. Decode with `decode_with_emissions(B_hat, ciphertext)` or feed emissions to Viterbi[8] for MAP decoding.

Log-likelihoods from a representative run are in the convergence table below.

EM convergence snapshot

Iteration	$\log P(V \Theta)$	Note
0	-374.15	random init
2	-297.81	structure starts to appear
16	-242.69	readable plaintext emerges
950	-237.72	plateau; restart to escape local optima

Scaling prevents underflow, and strictly positive A keeps the chain irreducible and aperiodic. Run multiple random restarts (different B seeds) and keep the run with the best final log-likelihood.

EM decoded sample

- Iter 0: BZN BQWNXLN JQBZ ... (random initialization, gibberish)
- Iter 16: THE TOXUPLE WITH HAZIND AN OQEN MING ... (largely readable, a few bad mappings)
- Iter 950: THE TROUPLE WITH HAZING AN OQEN MIND ... (plateau; minor letter swaps remain)

These snapshots show how likelihood climbs and the mapping improves before stalling. Rerun with new seeds to escape plateaus.

EM usage example: computing statistics for the English bigrams LM Use `build_bigram_model` to estimate the bigram transition matrix A and initial state distribution π from a cleaned English corpus (e.g., Brown). The helper trims the text to a configurable limit to keep runtimes reasonable.

EM usage: decrypt a text With A and π fixed, initialize an emission matrix B , apply `constrain_space_mapping`, and run `BaumWelch.process` with `recompute_A=False` and `recompute_PI=False`. Use `decode_with_emissions` for a quick readout or run Viterbi[8] with the learned emissions to obtain a MAP state sequence.

EM: Review of the Baum-Welch approach Baum-Welch can often recover readable plaintext quickly on sufficiently long ciphertexts, although runtime and quality depend on text length and initialization. In our setup we train A, π from an English corpus (bigram LM), then learn an emission matrix B per ciphertext while keeping A, π fixed.

Notes and limitations:

- the entropy of English implies a (rough) *unicity distance* - the minimum length needed for unique recovery in an idealized setting (Shannon, 1949) [9]. A common estimate uses per-character redundancy $D \approx 3.2$ bits:

$$U = \frac{H(K)}{D} = \frac{\log_2(26!)}{3.2} \approx 28$$

In practice, Baum-Welch typically needs substantially longer ciphertexts for stable convergence.

- Using only a first-order (bigram) language model limits context to local dependencies.
- Higher-order models (e.g., trigrams) can improve accuracy but increase computation and memory.
- A substitution cipher is a bijection, but plain Baum-Welch does not enforce a one-to-one constraint on B ; enforcing bijectivity requires constrained optimization or post-processing.
- We can bias initialization toward expected mappings (e.g., space→space, or other priors in B). Injecting higher-level knowledge (common words, names, greetings) is possible via constraints/priors or hybrid search, but it is not built into standard Baum-Welch.

Breaking the cipher with Metropolis-Hastings

Let's try a different angle.

Is there a way to enforce the bijective property of the cipher while solving the problem?

We briefly describe the Monte Carlo technique and the Metropolis-Hastings algorithm before using it with a new model.

Monte Carlo

If a quantity is intractable or hard to compute directly, we can approximate it via random sampling from a related distribution, avoiding the need for closed-form mathematics.

A typical example is estimating the value of π . Consider a circle of radius 1 inscribed in the square $[-1, 1] \times [-1, 1]$. If we repeatedly generate pseudo-random points (x, y) uniformly in the square and count the fraction that fall inside the circle ($x^2 + y^2 \leq 1$), that fraction estimates the area ratio:

$$\begin{aligned} \text{square area} &= 2r \cdot 2r = 4 \\ \text{circle area} &= \pi \cdot r^2 = \pi \\ \frac{\text{circle area}}{\text{square area}} &= \pi/4 \end{aligned}$$

Therefore $\pi \approx 4 \cdot (\text{fraction of points inside the circle})$. By the law of large numbers, this estimate converges as the number of samples grows.

```
1 limit=1_000_000 # the more the better
2 print(sum(random.uniform(-1,1)**2 + random.uniform(-1,1)**2<=1 for i in range(limit))
     )*4.0/limit)
3 3.14188
```

Under “Monte Carlo” methods fall different techniques, which stochastically tweak the parameters of a model to explore unknown values and evaluate its outputs, such that the model can be improved. In our setting, this “tweaking” will take the form of a random-walk proposal over candidate cipher keys. This walk is mostly done by generating random samples and then accepting or discarding them according to a criterion, or by generating samples from some weighted distribution (importance sampling). For example, in the case of a substitution cipher, consider a random walk where at each step two letters of the cipher key are swapped and the resulting plaintext is evaluated for plausibility under a language model.

Metropolis-Hastings, brief introduction

Metropolis-Hastings is a Markov chain Monte Carlo (MCMC) algorithm that works well for the substitution cipher case.

- consider a target distribution $P(x)$ over solutions x (here: cipher keys) which we cannot evaluate exactly, but for which we can compute a score from the observations;
- we have a function $\ell(x) = \log \tilde{P}(x) + C$ which is equal to $\log P(x)$ up to an additive constant (log unnormalized score);
- we make a random walk by randomly making a change to the parameters (e.g., swapping two letters in the key) and then evaluating the result;
- the new value is accepted or rejected using an acceptance function that compares $\tilde{P}(x')$ with $\tilde{P}(x)$ (and, in general, also accounts for the proposal distribution);
- the proposal distribution $q(x'|x)$ specifies how we generate the candidate x' from the current state x (for swap moves, it is typically symmetric).

In the longer term (and under some basic conditions), this walk is guaranteed to sample from the target distribution $P(x)$ (more precisely: $P(x)$ is the stationary distribution of the Markov chain).

More about convergence of Markov models

A discrete-time Markov chain with transition probabilities $P(y|x)$ has a stationary distribution Π^* if

$$\Pi^*(y) = \sum_x \Pi^*(x)P(y|x)$$

In words: if the chain starts distributed as Π^* , then after one step it is still distributed as Π^* .

If the chain is irreducible and aperiodic (and the state space is finite), then a stationary distribution exists, is unique, and the chain converges to it. In particular, if we write the distribution at time t as a row vector Π_t , and the transition matrix as A with entries $A_{xy} = P(y|x)$, then

$$\Pi_{t+1} = \Pi_t A \quad \Pi_t = \Pi_0 A^t$$

For large t ,

$$\Pi_0 A^t \xrightarrow[t \rightarrow \infty]{} \Pi^*$$

and the stationary distribution satisfies the fixed-point equation

$$\Pi^* = \Pi^* A$$

Note that convergence does not mean the chain stops moving; it means the distribution of the state stabilizes.

A strong sufficient condition is strict positivity: if $A_{xy} > 0$ for all x, y , then the chain is automatically irreducible and aperiodic, hence it has a unique stationary distribution Π^* .

In Metropolis-Hastings, the goal is to construct a Markov chain whose stationary distribution is the desired target distribution over solution (e.g., cipher keys). In other words, the target distribution is the Π^* we want the chain to converge to.

A common way to guarantee that a chosen Π^* is stationary is detailed balance (time reversibility):

$$\Pi^*(x)P(y|x) = \Pi^*(y)P(x|y) \tag{3}$$

Summing both sides over x gives stationarity:

$$\sum_x \Pi^*(x)P(y|x) = \sum_x \Pi^*(y)P(x|y) = \Pi^*(y) \sum_x P(x|y) = \Pi^*(y)$$

since $\sum_x P(x|y) = 1$.

Some clues behind the inner workings of Metropolis-Hastings for the substitution cipher

So far, we modeled ciphertext with an HMM and learned an emission model with Baum-Welch. However, the substitution cipher key is bijective, and Baum-Welch does not naturally enforce a one-to-one constraint on the emission matrix. Also, using higher-order language models inside an HMM quickly becomes expensive. Let's rethink the model in a way that enforces the substitution mapping directly, and makes it easy to plug in an n -gram LM.

The technique discussed here is inspired by “The Markov Chain Monte Carlo Revolution” (Diaconis, 2009) [10] which gives deeper insights into MCMC and uses Metropolis-Hastings to (among other things) decode an encrypted message.

Consider the substitution mapping S (a bijection) and its inverse S^{-1} : $S(P) = C$ and $S^{-1}(C) = P$.

In Metropolis-Hastings, the state of the Markov chain is the key S itself. We move from one key to another by proposing small changes (e.g., swapping two letters in the key) and accepting/rejecting them. The chain is constructed so that its stationary distribution $\Pi(S)$ assigns higher probability to keys that produce more plausible English plaintext.

The stationary distribution $\Pi(S)$ does not need to be known in normalized form: it is enough to define an unnormalized score $\tilde{\Pi}(S) \propto \Pi(S)$, because MH only uses ratios of scores.

Using an n -gram language model (LM), we can score a candidate key S by decoding the ciphertext sequence (c_1, \dots, c_N) into candidate plaintext symbols

$$p_i = S^{-1}(c_i)$$

and then computing the LM score of the resulting plaintext under the n -gram factorization:

$$Pl(S) = \sum_{i=n}^N \log P_{LM}(p_i \mid p_{i-n+1:i-1})$$

In implementation we use an unnormalized weighted sum of n-gram log-weights; MH only needs score differences.

Metropolis-Hastings (symmetric proposal): the algorithm

Metropolis et al. (1953) [11] is the special case of Metropolis-Hastings where the proposal is symmetric; it can be motivated via detailed balance.

It needs (i) a symmetric proposal distribution to generate candidate moves and (ii) a plausibility score (based on the language model) proportional to the target distribution. It works as follows:

0. Start with:

- a current state x (e.g., a key; can be random),
- a symmetric proposal $g(y|x)$ such that $g(y|x) = g(x|y)$,
- a log-score $Pl(x)$ (plausibility), defined so that $\exp(Pl(x)) \propto \Pi(x)$ (equivalently, $Pl(x) = \log \tilde{\Pi}(x) + C$ for some constant C).

1. From the current state x , draw a candidate $y \sim g(y|x)$.
2. Compute the acceptance probability

$$\alpha(x, y) = \min(1, \exp(Pl(y) - Pl(x)))$$

3. Accept/Reject:

- draw $u \sim \text{Uniform}(0, 1)$;
- if $u \leq \alpha(x, y)$, accept and set $x \leftarrow y$;
- otherwise, reject and keep x .

4. Repeat from step 1 (for a fixed number of iterations, or after burn-in collect samples).

The acceptance rule comes from the detailed-balance condition (3). Let Π^* be the desired stationary distribution and let $T(y|x)$ be the actual transition kernel of the Metropolis chain. Condition (3) requires:

$$\Pi^*(x)T(y|x) = \Pi^*(y)T(x|y)$$

For $y \neq x$, the Metropolis transition can be written as:

$$T(y|x) = g(y|x)\alpha(x,y)$$

where $g(y|x)$ is the proposal distribution and $\alpha(x,y)$ is the acceptance probability. Plugging this into (3) gives:

$$\frac{g(y|x)\alpha(x,y)}{g(x|y)\alpha(y,x)} = \frac{\Pi^*(y)}{\Pi^*(x)}$$

If the proposal is symmetric, $g(y|x) = g(x|y)$, this simplifies to:

$$\frac{\alpha(x,y)}{\alpha(y,x)} = \frac{\Pi^*(y)}{\Pi^*(x)}$$

Since $Pl(x)$ is a log-score (i.e., $Pl(x) = \log \tilde{\Pi}(x)$ up to an additive constant), the condition is satisfied by choosing:

$$\alpha(x,y) = \min(1, \exp(Pl(y) - Pl(x))) = \min\left(1, \frac{\Pi^*(y)}{\Pi^*(x)}\right)$$

```

1 ct = normalize(ciphertext)
2 S = random_key(alphabet)                      # current bijection (state)
3 pt = decode(ct, S)                            # candidate plaintext under S
4 score = plausibility(pt, lm)                  # log-score (recommended)
5
6 best_S, best_score = S, score
7
8 for _ in range(N):
9     a, b = pick_two_symbols(alphabet)
10    S2 = swap_in_key(S, a, b)                  # symmetric proposal, still bijective
11    pt2 = decode(ct, S2)
12    score2 = plausibility(pt2, lm)Δ
13
14    = score2 - score                         # log acceptance ratio
15    if log(rand_uniform()) <= Δ:           # accept with prob min(1, expΔ())
16        S, pt, score = S2, pt2, score2
17
18    if score > best_score:
19        best_S, best_score = S, score

```

The proposal enforces a bijection by swapping two symbols; the plausibility function can mix multiple n-gram lengths.

Acceptance intuition: early on, many swaps are accepted; later, acceptance drops as the chain reaches higher-scoring regions. Restart from a fresh seed if trapped in low-quality states.

Plausibility should return a log score (sum of log n-gram probs) so Δ is stable and $\exp(\Delta)$ is well-defined.

Metropolis-Hastings implementation: a concrete example

With this algorithm, we can exploit an LM with multiple n -grams. This is straightforward because the plausibility function only needs to rank candidates; a fully normalized probability is not required. Intuitively, a candidate plaintext should score higher when it contains n -grams that are typical for English. A simple choice is the sum of weighted n -gram frequencies, where the weight increases with n because longer contexts carry more information.

Since “frequency” is a normalized quantity, the normalization step can be omitted as long as the normalization factor (e.g., the total number of extracted n -grams) is constant across candidates; it cancels out when comparing two states in the acceptance ratio. In practice, the weighted count is treated as a log-plausibility score, so the acceptance step uses $\exp(\Delta)$ with $\Delta = Pl(y) - Pl(x)$.

Furthermore, at each step an existing mapping is swapped; this maintains the bijection between the plaintext and ciphertext alphabets.

The n -gram model is sparse for larger n , so it would be inefficient to store it as a dense array. In this Python example we use a dictionary to store only the observed n -grams.

Finally, the acceptance rule allows occasional moves to lower-scoring states, which helps the chain jump between local optima. There is no single point where the algorithm is guaranteed to be “done”, so for this project we set a maximum number of iterations; another option is to stop if there is no improvement in the best score for a fixed window of iterations.

Here is my full implementation of the Metropolis algorithm and a concrete example:

MH usage sketch (see Appendix B for full code)

- Build a multi-gram LM: `lm = count_ngrams(clean_corpus, min_len=2, max_len=5)`.
- Substitution cipher: `metropolis_substitution(ciphertext, lm, iterations=50k, seed=7)` tracks and prints intermediate best candidates during the run.
- Patristocrat spacing: `metropolis_patristocrat(text_without_spaces, lm, iterations=5k, seed=11)` proposes insert/remove-space moves and tracks the best spacing found.

MH trace highlights

- In this trace, the first accepted swaps quickly lock in a few high-signal bigrams (e.g., TH, HE), and readability starts improving early.
- The acceptance rate here tracks the plausibility spread: once the score separation between candidates grows, most proposed swaps are rejected and progress comes from occasional improving moves.
- In the spacing stage of this same trace, patristocrat spacing recovers as common trigrams reappear.

MH trace sample

- It#1: K CXJO XI HBD BKIO XS ENJHB HEN XI HBD CLSB (initial guess)
- It#15: M CFJS FI TBR BMIS FO ENJTB TEN FI TBR CLOB (bigrams emerging)
- It#219: IF YOU LOOK FOR PERFECTION YOU LL NEVER BE CONTENT (spacing recovered)

Patristocrats

To make a substitution cipher harder to break, spaces can be removed from the ciphertext. The American Cryptogram Association calls this kind of cryptogram a “Patristocrat”. This increases difficulty, at the cost of added ambiguity when reconstructing word boundaries.

To solve a Patristocrat, it is convenient to score candidate plaintext without spaces using an LM trained on a corpus with spaces removed (alphabet A-Z only). The same algorithms as before can be used, but longer ciphertext is typically needed to provide enough statistical signal. Higher-order n -gram language models (as

used in the Metropolis-Hastings scoring function) tend to work much better here than first-order/bigram models, as expected.

After the letters have been decoded, spaces can be reintroduced using an LM trained on a corpus with spaces. There are several ways to do this (e.g., dynamic programming), but with only minor changes to the previous implementation we use Metropolis-Hastings: proposals insert/remove a space at random positions and moves are accepted according to the spaced-LM plausibility score.

Practical usage: train the A-Z LM (spaces removed) and run `metropolis_patristocrat`; then train the spaced LM and run the spacing sampler.

Review of the Metropolis approach

With enough iterations and a few random restarts, Metropolis-Hastings can often recover readable plaintext even from relatively short ciphertexts, although success depends on text length and the strength of the language-model score. The algorithm is flexible: it can incorporate higher-order n-gram language models directly in the plausibility function without changing the sampler itself. In practice, the speed of improvement depends mainly on the proposal function (how new candidates are generated) and the acceptance rule (which controls how often worse moves are tolerated). Variants of Metropolis-Hastings typically modify these two components to trade off exploration, runtime, and solution quality.

Limitations & future work

- Short ciphertexts remain ambiguous; try multiple random seeds/restarts or incorporate partial known-plaintext cribs.
- Higher-order n -gram LMs can improve plausibility, but they require smoothing (to avoid zeros) and larger corpora to reduce sparsity.
- Other languages and alphabets (diacritics, punctuation, different whitespace rules) require retraining the LM and adjusting the symbol mapping/preprocessing.
- EM can stall in local optima; multiple restarts can help, and annealed/tempered variants of EM may improve robustness.

Appendices

Full Python listings referenced in the text. Place these in `src/crypto/` or on your import path when re-running the experiments.

Appendix A: EM helper code

```
1 """
2 Baum-Welch routines tailored to simple substitution ciphers.
3 """
4
5 from __future__ import annotations
6
7 from typing import Callable, Tuple
8
9 import numpy as np
10
11 from .utils import (
12     ALPHABET,
13     SPACE_IDX,
14     char_to_index,
15     encode_text,
16     normalize_text,
```

```

17     seed_everything,
18 )
19
20 EmissionHook = Callable[[np.ndarray], np.ndarray]
21
22
23 class BaumWelch:
24     """Minimal Baum-Welch implementation with scaling to avoid underflow."""
25
26     def __init__(self, num_hidden: int = len(ALPHABET)):
27         self.num_hidden = num_hidden
28         self.num_visible = num_hidden
29         self.T = 0
30
31     @staticmethod
32     def normalize(matrix: np.ndarray) -> Tuple[np.ndarray, np.ndarray]:
33         totals = np.sum(matrix, axis=min(1, len(matrix.shape) - 1), keepdims=True)
34         totals[totals == 0] = 1.0
35         return matrix / totals
36
37     def alpha(self, A: np.ndarray, B: np.ndarray, PI: np.ndarray, V: np.ndarray):
38         alphas = np.zeros((self.T, self.num_hidden))
39         scale_facts = np.zeros(self.T)
40         alphas[0], scale_facts[0] = self.normalize(PI * B[:, V[0]])
41         for t in range(1, self.T):
42             alphas[t], scale_facts[t] = self.normalize(
43                 B[:, V[t]] * np.sum(alphas[t - 1] * A.T, axis=1)
44             )
45         return alphas, scale_facts
46
47     def beta(self, A: np.ndarray, B: np.ndarray, V: np.ndarray, scale_facts: np.ndarray):
48         betas = np.zeros((self.T, self.num_hidden))
49         betas[self.T - 1] = np.array([1.0 / scale_facts[self.T - 1]])
50         for t in range(self.T - 2, -1, -1):
51             betas[t] = (
52                 np.sum(betas[t + 1] * B[:, V[t + 1]] * A, axis=1) / scale_facts[t]
53             )
54         return betas
55
56     def process(
57         self,
58         A: np.ndarray,
59         B: np.ndarray,
60         PI: np.ndarray,
61         V: np.ndarray,
62         max_iter: int,
63         recompute_A: bool = True,
64         recompute_B: bool = True,
65         recompute_PI: bool = True,
66         ext_fun: EmissionHook | None = None,
67         verbose_fun: Callable[[np.ndarray, np.ndarray, int, float], None] | None = None,
68     ):
69         self.T = V.size
70         log_scale_facts = None

```

```

71     for it in range(max_iter):
72         alphas, scale_facts = self.alpha(A, B, PI, V)
73         new_log_scale = float(np.sum(np.log(scale_facts)))
74         if log_scale_facts is not None and new_log_scale <= log_scale_facts:
75             if verbose_fun:
76                 print("Algorithm has converged")
77             break
78         log_scale_facts = new_log_scale
79         betas = self.beta(A, B, V, scale_facts)
80         gammas = np.zeros((self.T, self.num_hidden, self.num_hidden))
81         for t in range(self.T - 1):
82             gammas[t] = (A.T * alphas[t, :]).T * (betas[t + 1] * B[:, V[t + 1]])
83             gammas_i = np.sum(gammas, axis=2)
84             gammas_i[self.T - 1] = alphas[self.T - 1, :]
85
86         if recompute_A:
87             gammas_t = np.sum(gammas, axis=0)
88             totals = np.sum(gammas_t, axis=1)[:, None]
89             totals[totals == 0] = 1.0
90             A = (gammas_t.T / totals).T
91
92         if recompute_B:
93             for v in range(self.num_visible):
94                 mask = V == v
95                 if np.any(mask):
96                     B[:, v] = np.sum(gammas_i[mask], axis=0)
97             totals = np.sum(gammas_i, axis=0)
98             totals[totals == 0] = 1.0
99             B = (B.T / totals).T
100
101        if recompute_PI:
102            PI = gammas_i[0]
103
104        if ext_fun is not None:
105            B = ext_fun(B)
106
107        if verbose_fun:
108            verbose_fun(V, B, it, log_scale_facts)
109
110    return A, B, PI, log_scale_facts
111
112
113 def constrain_space_mapping(B: np.ndarray) -> np.ndarray:
114     """
115     Force a 1-1 mapping for space:space in the emission matrix.
116
117     Ensures space can only map to space, keeping the substitution mapping bijective for that
118     symbol.
119     """
120     B = B.copy()
121     B[SPACE_IDX, :] = 0.0
122     B[:, SPACE_IDX] = 0.0
123     B[SPACE_IDX, SPACE_IDX] = 1.0
124     row_sums = B.sum(axis=1, keepdims=True)

```

```

124     row_sums[row_sums == 0] = 1.0
125     return B / row_sums
126
127
128 def build_bigram_model(
129     corpus_text: str,
130     *,
131     max_iter: int = 150,
132     limit: int = 1_000_000,
133     seed: int = 13,
134 ):
135     """
136     Train a bigram language model from plain text using Baum-Welch.
137
138     Returns the transition matrix A, an identity emission matrix B, and the
139     initial state distribution PI. The text is upper-cased, cleaned, and
140     truncated to `limit` characters to keep runtimes reasonable.
141     """
142     seed_everything(seed)
143     normalized = normalize_text(corpus_text)[:limit]
144     V = np.array(encode_text(normalized))
145
146     num_states = len(ALPHABET)
147     PI = np.random.rand(num_states)
148     A = np.random.rand(num_states, num_states)
149     B = np.identity(num_states)
150
151     baum = BaumWelch(num_states)
152     A, _, PI, _ = baum.process(A, B, PI, V, max_iter=max_iter, recompute_B=False)
153     return A, B, PI
154
155
156 def decode_with_emissions(B: np.ndarray, observations: str) -> str:
157     """
158     Simple decoder: pick the highest emission probability per symbol.
159
160     For maximum a posteriori decoding, use the Viterbi algorithm instead.
161     """
162     V = [char_to_index(ch) for ch in observations]
163     decoded_indices = [int(np.argmax(B[:, v])) for v in V]
164     return ''.join(ALPHABET[idx] for idx in decoded_indices)

```

Appendix B: Metropolis-Hastings helper code

```

1 """
2 Metropolis-Hastings helpers for substitution and patristocrat variants.
3 """
4
5 from __future__ import annotations
6
7 import math
8 import random
9 import string
10 from collections import Counter

```

```

11 from typing import Dict, Iterable
12
13 from .utils import normalize_text
14
15
16 def count_ngrams(text: str, min_len: int = 2, max_len: int = 9) -> Dict[int, Dict[str, float]]:
17     """Count n-grams of varying length and return log counts for stability."""
18     return {
19         n: {k: math.log(v) for k, v in Counter(text[idx : idx + n] for idx in range(len(text)) - n + 1).items()}
20         for n in range(min_len, max_len + 1)
21     }
22
23
24 def plausibility_score(text: str, lm: Dict[int, Dict[str, float]]) -> float:
25     """Score text against an LM, weighted by n-gram length."""
26     counts = count_ngrams(text, min(lm), max(lm))
27     return sum(
28         lm[n].get(k, 0.0) * n
29         for n in counts
30         for k, _ in counts[n].items()
31         if k in lm[n]
32     )
33
34
35 def metropolis_substitution(
36     ciphertext: str,
37     lm: Dict[int, Dict[str, float]],
38     *,
39     iterations: int = 1_000_000,
40     seed: int = 13,
41 ) -> str:
42     """
43     Symmetric-proposal Metropolis-Hastings for simple substitution ciphers.
44
45     Fixed symbols (anything outside A-Z) are kept in place. The function yields
46     intermediate improvements to make convergence observable in static logs.
47     """
48     random.seed(seed)
49     ciphertext = normalize_text(ciphertext)
50     best = current_score = plausibility_score(ciphertext, lm)
51
52     fixed = set(ciphertext) - set(string.ascii_uppercase)
53     used_symbols = list(set(ciphertext) - fixed)
54     all_symbols = set(string.ascii_uppercase)
55     accepted = ciphertext
56     for i in range(iterations):
57         c1 = random.choice(used_symbols)
58         c2 = random.choice([c for c in all_symbols if c != c1])
59         candidate = accepted.translate(str.maketrans({c1: c2, c2: c1}))
60         candidate_score = plausibility_score(candidate, lm)
61         mu = math.log(random.uniform(0, 1))
62         if mu < max(min(candidate_score - current_score, 0), math.log(1e-3)):

```

```

63         accepted = candidate
64         current_score = candidate_score
65         used_symbols = list(set(accepted) - fixed)
66         if candidate_score > best:
67             best = candidate_score
68             print(f"It#{i}: {accepted}")
69     return accepted
70
71
72 def metropolis_patristocrat(
73     ciphertext: str,
74     lm: Dict[int, Dict[str, float]],
75     *,
76     iterations: int = 5000,
77     seed: int = 21,
78 ) -> str:
79     """
80     Restore spaces in patristocrat-style ciphertexts using a MH sampler.
81
82     The proposal either removes a space or inserts one adjacent to letters. A
83     length-normalized plausibility score keeps proposals comparable.
84     """
85     random.seed(seed)
86     ciphertext = normalize_text(ciphertext).replace(" ", "")
87     best = current = ciphertext
88     best_score = current_score = plausibility_score(ciphertext, lm) / len(ciphertext)
89
90     for i in range(iterations):
91         c = random.choice(range(1, len(current) - 1))
92         candidate = list(current)
93         if current[c].isspace():
94             candidate.pop(c)
95         elif current[c - 1] != " " and current[c + 1] != " ":
96             candidate.insert(c, " ")
97         candidate_text = "".join(candidate)
98         cand_score = plausibility_score(candidate_text, lm) / len(candidate_text)
99         mu = math.log(random.uniform(0, 1))
100        if mu < min(cand_score - current_score, 0):
101            current = candidate_text
102            current_score = cand_score
103            if cand_score > best_score:
104                best = current
105                best_score = cand_score
106                print(best)
107    return best

```

Appendix C: Utility functions

```

1 import random
2 import re
3 from typing import Iterable
4
5 ALPHABET = "ABCDEFGHIJKLMNOPQRSTUVWXYZ "
6 SPACE_IDX = ALPHABET.index(" ")

```

```

7
8
9 def seed_everything(seed: int) -> None:
10     """Seed Python's random generator and numpy if available."""
11     random.seed(seed)
12     try:
13         import numpy as np
14
15         np.random.seed(seed)
16     except Exception:
17         # numpy is optional
18         pass
19
20
21 def normalize_text(text: str) -> str:
22     """
23     Keep A-Z and spaces, collapse whitespace, and uppercase.
24
25     Ensures corpus and cipher text share the same character mapping for the LM.
26     """
27     cleaned = re.sub("[^A-Z ]", " ", text.upper())
28     return re.sub("\s+", " ", cleaned).strip()
29
30
31 def char_to_index(ch: str) -> int:
32     """Map a character in ALPHABET to its ordinal index."""
33     return ALPHABET.index(ch)
34
35
36 def index_to_char(idx: int) -> str:
37     """Map an ordinal index back to a character in ALPHABET."""
38     return ALPHABET[idx]
39
40
41 def encode_text(text: Iterable[str]) -> list[int]:
42     """Convert a string of characters into indices."""
43     return [char_to_index(ch) for ch in text]
44
45
46 def decode_indices(indices: Iterable[int]) -> str:
47     """Convert indices back to a string of characters."""
48     return "".join(index_to_char(i) for i in indices)

```

Appendix D: Package exports

```

1 """
2 Lightweight cryptanalytic helpers for EM- and MH-based substitution cipher exploration.
3 """
4
5 from .utils import (
6     ALPHABET,
7     SPACE_IDX,
8     char_to_index,
9     index_to_char,

```

```

10     encode_text,
11     decode_indices,
12     normalize_text,
13     seed_everything,
14 )
15 from .em import BaumWelch, build_bigram_model, constrain_space_mapping
16 from .mh import (
17     count_ngrams,
18     plausibility_score,
19     metropolis_substitution,
20     metropolis_patristocrat,
21 )
22
23 __all__ = [
24     "ALPHABET",
25     "SPACE_IDX",
26     "char_to_index",
27     "index_to_char",
28     "encode_text",
29     "decode_indices",
30     "normalize_text",
31     "seed_everything",
32     "BaumWelch",
33     "build_bigram_model",
34     "constrain_space_mapping",
35     "count_ngrams",
36     "plausibility_score",
37     "metropolis_substitution",
38     "metropolis_patristocrat",
39 ]

```

Appendix E: Demo entrypoint

```

1 """Command-line entrypoint for the Baum-Welch and Metropolis demos."""
2
3 from __future__ import annotations
4
5 import argparse
6 import math
7 from pathlib import Path
8
9 import numpy as np
10
11 from .em import BaumWelch, constrain_space_mapping, decode_with_emissions
12 from .mh import load_pickled_language_model, metropolis_patristocrat, metropolis_substitution
13 from .utils import ALPHABET, encode_text
14
15
16 def run_em_demo(max_iter: int = 500, verbose: bool = False, restarts: int = 15) -> None:
17     """Reproduce the Baum-Welch demo using the pickled bigram LM (no corpus training)."""
18     num_states = 27
19     lm_path = Path("resources/en/ngrams_space.pkl")
20     lm_stats = load_pickled_language_model(lm_path)
21

```

```

22     def _softmax_from_logs(logs: np.ndarray) -> np.ndarray:
23         m = float(np.max(logs))
24         shifted = logs - m
25         w = np.exp(shifted)
26         s = float(np.sum(w))
27         if s == 0.0 or not np.isfinite(s):
28             return np.full_like(w, 1.0 / len(w))
29         return w / s
30
31     def _build_bigram_language_model_from_pickle():
32         symbols = list(ALPHABET)
33         log_eps = math.log(1e-6)
34         log_pi = np.array([float(lm_stats[1].get(ch, log_eps)) for ch in symbols],
35                           dtype=float)
35         PI_local = _softmax_from_logs(log_pi)
36         A_local = np.zeros((len(symbols), len(symbols)), dtype=float)
37         for i, a in enumerate(symbols):
38             row_logs = np.array([float(lm_stats[2].get(a + b, log_eps)) for b in symbols],
39                               dtype=float)
40             A_local[i, :] = _softmax_from_logs(row_logs)
41         return A_local, PI_local
42
43     A, PI = _build_bigram_language_model_from_pickle()
44
45     # use Baum-Welch to decrypt this text
46     encrypted = (
47         "RBO RPKTIGO VCRB BWUCJA WJ KLOJ HCJD KM SKTPQO CQ RBWR "
48         "LOKLGO VCGG CJQCQR KJ SKHCJA WGKJA WJD RPYCJA RK LTR RBCJAQ CJ CR "
49     )
50     V = np.array(encode_text(encrypted))
51     baum = BaumWelch()
52
53     log_eps = math.log(1e-6)
54     bigram = lm_stats[2]
55
56     def score_bigram(text: str) -> float:
57         # sum bigram log-weights; higher is better
58         total = 0.0
59         for i in range(len(text) - 1):
60             bg = text[i : i + 2]
61             total += bigram.get(bg, log_eps)
62         return total
63
64     best_overall = {"text": "", "log": -1e300, "score": -1e300}
65
66     for r in range(restarts):
67         # Start from a random emission matrix; space forced later.
68         B = np.random.rand(num_states, num_states)
69
70         best_candidate = {"text": "", "log": -1e300, "score": -1e300, "B": None}
71
72         def verbose_fun(V_, B_, it, log_norm):
73             decoded = decode_with_emissions(B_, encrypted)
74             sc = score_bigram(decoded)

```

```

74     if verbose:
75         print(f"[r{r}] #It", it, decoded, log_norm, f"(bg_score={sc:.2f})")
76     if sc > best_candidate["score"]:
77         best_candidate["score"] = sc
78         best_candidate["text"] = decoded
79         best_candidate["log"] = log_norm
80         best_candidate["B"] = B_.copy()
81     if log_norm > best_candidate["log"]:
82         best_candidate["log"] = log_norm
83
84     _, B_hat, _, log_norm = baum.process(
85         A,
86         B,
87         PI,
88         V,
89         max_iter=max_iter,
90         ext_fun=constrain_space_mapping,
91         verbose_fun=verbose_fun,
92         recompute_A=False,
93         recompute_PI=False,
94         stop_on_decrease=False,
95     )
96     winner_decoded = best_candidate["text"] or decode_with_emissions(B_hat, encrypted)
97     winner_log = best_candidate["log"] if best_candidate["text"] else log_norm
98     winner_score = best_candidate["score"] if best_candidate["text"] else
99         score_bigram(winner_decoded)
100    # Prefer the emissions matrix at the best LM score if captured
101    if best_candidate["B"] is not None and best_candidate["text"]:
102        winner_decoded = decode_with_emissions(best_candidate["B"], encrypted)
103        if winner_score > best_overall["score"]:
104            best_overall = {"text": winner_decoded, "log": winner_log, "score": winner_score}
105
106    print(best_overall["text"])
107    print("log_norm:", best_overall["log"])
108
109 def run_mh_substitution_demo(iterations: int = 1_000_000) -> None:
110     """Metropolis substitution demo."""
111     print("Metropolis substitution:")
112     best = metropolis_substitution(
113         " K CXJO XI HBD BKIO XL ENJHB HEN XI HBD CSLB ",
114         iterations=iterations,
115         lm_pickle=Path("resources/en/ngrams_space.pkl"),
116     )
117     print("\nBest candidate:")
118     print(best)
119
120
121 def run_mh_spacing_demo(iterations: int = 5000) -> None:
122     """Metropolis spacing demo."""
123     print("\nMetropolis spacing:")
124     spaced, score = metropolis_patristocrat(
125         " IFYOULOOKFORPERFECTIONYOUULLNEVERBECONTENT ",
126         iterations=iterations,

```

```

127     lm_pickle=Path("resources/en/ngrams_space.pkl"),
128 )
129 print("\nBest spacing:")
130 print(spaced)
131
132
133 def main() -> int:
134     parser = argparse.ArgumentParser(description="Run the EM and MH demos from the examples.")
135     sub = parser.add_subparsers(dest="cmd", required=True)
136
137     p_em = sub.add_parser("em", help="Run Baum-Welch demo")
138     p_em.add_argument("--max-iter", type=int, default=5000)
139     p_em.add_argument("--verbose", action="store_true")
140     p_em.add_argument("--restarts", type=int, default=15)
141
142     p_mh = sub.add_parser("mh", help="Run Metropolis demos")
143     p_mh.add_argument("--iters-sub", type=int, default=1_000_000)
144     p_mh.add_argument("--iters-space", type=int, default=5000)
145
146     p_all = sub.add_parser("all", help="Run both demos")
147     p_all.add_argument("--max-iter", type=int, default=5000)
148     p_all.add_argument("--verbose", action="store_true")
149     p_all.add_argument("--iters-sub", type=int, default=1_000_000)
150     p_all.add_argument("--iters-space", type=int, default=5000)
151
152     args = parser.parse_args()
153
154     if args.cmd in {"em", "all"}:
155         run_em_demo(max_iter=args.max_iter, verbose=args.verbose, restarts=args.restarts)
156     if args.cmd in {"mh", "all"}:
157         run_mh_substitution_demo(iterations=args.iters_sub)
158         run_mh_spacing_demo(iterations=args.iters_space)
159
160
161
162 if __name__ == "__main__":
163     raise SystemExit(main())

```

References

- [1] Substitution cipher. (2021, July 27). In Wikipedia.
https://en.wikipedia.org/w/index.php?title=Substitution_cipher&oldid=1035776944
- [2] Dempster, A.P.; Laird, N.M.; Rubin, D.B. (1977). “Maximum Likelihood from Incomplete Data via the EM Algorithm”. Journal of the Royal Statistical Society, Series B. 39 (1): 1-38. JSTOR 2984875. MR 0501537.
- [3] Maya R. Gupta and Yihua Chen (2011), “Theory and Use of the EM Algorithm”, Foundations and Trends® in Signal Processing: Vol. 4: No. 3, pp 223-296. <https://dx.doi.org/10.1561/2000000034>
- [4] Knight, Kevin & Nair, Anish & Rathod, Nishit & Yamada, Kenji. (2006). Unsupervised Analysis for Decipherment Problems. <https://dx.doi.org/10.3115/1273073.1273138>
- [5] Gagniuc, Paul A. (2017). Markov Chains: From Theory to Implementation and Experimentation. USA, NJ: John Wiley & Sons. pp. 1-256. ISBN 978-1-119-38755-8

- [6] L. R. Rabiner, "A tutorial on hidden Markov models and selected applications in speech recognition," in Proceedings of the IEEE, vol. 77, no. 2, pp. 257-286, Feb. 1989. <https://dx.doi.org/10.1109/5.18626>
- [7] Francis, W. Nelson & Henry Kucera. 1967. Computational Analysis of Present-Day American English. Providence, RI: Brown University Press.
- [8] A. Viterbi, "Error bounds for convolutional codes and an asymptotically optimum decoding algorithm," in IEEE Transactions on Information Theory, vol. 13, no. 2, pp. 260-269, April 1967, doi: 10.1109/TIT.1967.1054010.
- [9] Shannon, Claude. "Communication Theory of Secrecy Systems", Bell System Technical Journal, vol. 28(4), page 656-715, 1949. <https://doi.org/10.1002/j.1538-7305.1949.tb00928.x>
- [10] Diaconis, Persi. (2009). The Markov Chain Monte Carlo Revolution. Bulletin of the American Mathematical Society. 46. 179-205. <https://dx.doi.org/10.1090/S0273-0979-08-01238-X>
- [11] Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., & Teller, E. (1953). Equation of State Calculations by Fast Computing Machines. The Journal of Chemical Physics, 21(6), 1087–1092. <https://doi.org/10.1063/1.1699114>