

Breaking substitution ciphers with Markov models

Algorithms description and implementation notes

Enrico Rollando (e.rollando@outlook.com)

August 6, 2021

Abstract

Cryptography can be an interesting field to practice the knowledge of statistical methods. In this document I use the simple substitution cipher as a toy example to decrypt messages without knowing the key. I focus on unsupervised learning techniques based on Markov models, review the theory behind them, and show concrete implementations and usage examples.

Goal & Scope

- Ciphertext alphabet: uppercase A–Z plus space. Spaces are preserved for the base examples; the patristocrat variant removes them.
- Corpus: Brown corpus (upper-cased, non-letters stripped, whitespace collapsed); any cleaned English corpus can be swapped in.
- Success criteria: recover readable plaintext and a near-bijective substitution map; report convergence traces instead of only final text.
- Deliverables: pseudocode for EM and Metropolis-Hastings, runnable helpers in `src/crypto`, and reproducible setup (seed + package versions).

How to reproduce

- Inputs: cleaned English corpus (Brown recommended), ciphertext samples in A–Z + space, patristocrat sample without spaces.
- Runtime (typical laptop): EM bigram LM ~2–4 minutes at 200k chars, EM decryption ~1–2 minutes at 200 iterations; MH 50k swaps on short texts ~seconds; patristocrat spacing ~seconds.
- Seed: fix RNG seeds (e.g., 13/42) for comparable traces; exact ciphertext-dependent results still vary.

Reproducibility setup

- Environment: Python 3.x, `numpy`, `nltk` (for Brown corpus) or any cleaned English corpus.
- Seeding: fix RNG seeds for repeatable EM/MH traces.
- Implementation: place the Python modules from Appendices A–C on your `PYTHONPATH` or in a local `src/crypto` package.
- Corpus: upper-case, strip non-letters (except space), collapse whitespace.

All subsequent code references point to the appendices rather than to executable cells, since this document targets PDF publication.

Corpus note: examples assume an English corpus such as NLTK’s Brown (`nltk.corpus.brown`)[7]. For offline or non-NLTK use, substitute any cleaned English corpus (A–Z plus space, whitespace collapsed).

Introduction

Can we break a simple substitution cipher automatically? “To break” here means: discover and use some weaknesses to restore - at least in part - the original text.

To define the scope, let's look only at plaintext and ciphertext composed in English capital letters and space, 26+1 symbols in total. Let's further say that the space is encrypted as space.

The *simple substitution cipher* is a mapping function $S_k(\text{text}) \rightarrow \text{ciphertext}$ where k is the key. It's a bijection between the alphabet sequence A and one of its $|A|!$ permutations. Given the ciphertext only, I want to figure out the inverse mapping $S^{-1}(\text{ciphertext})$ in an unsupervised way.

Let's try to define a hypothetical function: $\text{fit}(\text{text})$ that evaluates the current guess, ideally as "what is the probability that this text is the original plaintext?". A slightly different question that leads to a similar result and which I *can* actually define is: "how likely is it that this text is (English) plaintext?". I can use this function to evaluate and compare different guesses and choose the best one.

To model this evaluation I could use for example the unique features of a specific language. How often does a certain letter, or a certain group of letters (n-grams), appear in the plaintext. Or: how often a certain letter happens to follow another. And then look at the ciphertext and measure the distance using these statistics.

There is still a problem, the number of possible keys is huge (26! for the English alphabet), too many to try them all in a reasonable amount of time. Other techniques have to be used, in this document I show some of them.

Let us first look at the cipher to break and briefly present the Language Model which is the fundamental feature used to exploit the weakness of the cipher.

The simple substitution cipher

A simple substitution cipher [1] is a set of functions each modeling a one-to-one mapping between two alphabets: the plaintext- and the ciphertext- alphabet. This mapping is the key of the cipher.

Spaces are sometimes removed from plaintext before encryption, as they provide information about word boundaries, which in turn helps recovering the plaintext.

$$\begin{aligned} P &= [a_1, \dots, a_{26}], \text{example: ABCDEFGHIJKLMNOPQRSTUVWXYZ} \\ C &= [c_1, \dots, c_{26}], \text{example: HQUVNCFYJPXJDIMOBWLAZTGSRKE} \\ S: f(P) &\rightarrow C \\ S(\text{"THISISATEST"}) &= \text{"ZYPAPAHZNAZ"} \\ S^{-1}(\text{"ZYPAPAHZNAZ"}) &= \text{"THISISATEST"} \end{aligned} \tag{1}$$

The Language Model

In any language there are implicit (not formalized) rules that we use when we make or interpret a sentence. For example, an incomplete sentence "I like to write my letters by ..." would be expected to end with the word "hand". The initial words give a specific *context* and a context gives some words more *plausibility* than others.

At the character level, some knowledge of English is enough to read an incomplete word like "uncomfrt.b.y". The plausibility can be interpreted statistically: what is the most likely character in this position, given the visible sequence of characters?

The Language Model is tied to a specific language and focuses on frequencies of elements, on character or word level. As an example I consider the bigram English Language Model LM_{EN} , informally:

- bigram frequency:
 - how frequent is a bigram ab in English text? $P_{LM}(ab) = \frac{\text{count}(\text{occurrences of } ab)}{\text{len}(\text{text}) - 1}$
 - For a bigram ab how frequently does b follow a in English text? $P_{LM}(b|a) = \frac{\text{count}(ab)}{\text{count}(a)}$

The weakness

The simple substitution cipher encrypts the same symbol each time the same way. The ciphertext positionally maintains statistical information of the original plaintext. In other words, the i-th n-gram of both plaintext

and ciphertext will have the same frequency in its message. For example, if the second bigram is a double-letter, the most probable plaintext bigram is “TT” or “LL”.

To exploit this fact I could for example define a generative distribution G for the plaintext, controlled by some parameters Θ , then generate candidate plaintext, such that its positional n-grams distribution probability matches the one of the ciphertext.

This process could be progressively refined by automatically tweaking the parameters Θ to let distribution G converge to the observed distribution of the ciphertext.

Results at a glance

| Method | Cipher length | Iterations/time | Outcome |
|-------------------|---------------|-----------------------------|--|
| EM (Baum–Welch) | ~120 chars | 200 iters (~1–2 min) | Readable plaintext; some letters may need manual swap |
| MH (substitution) | ~40 chars | 50k swaps (~seconds) | Gradual convergence to legible text; prints best candidates |
| MH (patristocrat) | ~30 chars | 5k insert/remove (~seconds) | Recovers likely spacing; ambiguity remains on very short texts |
| Bigram LM prep | 200k chars | ~2–4 min | Transition matrix A, start dist PI reused across runs |

```

1 # scaled forward/backward to avoid underflow
2 initialize A, B, PI
3 repeat until convergence:
4   alpha, scale = forward(A, B, PI, V)
5   beta = backward(A, B, V, scale)
6   gamma_state = expected_state_visits(alpha, beta)
7   gamma_trans = expected_state_transitions(alpha, beta, A, B, V)
8   if recompute_A: A = normalize_rows(gamma_trans)
9   if recompute_B: B = normalize_rows(state_emissions(gamma_state, V))
10  if recompute_PI: PI = gamma_state[0]
11  B = constrain_space(B)  # keep → spacespace bijective

```

This highlights the scaled passes, the non-zero constraints on A , and the bijective tweak for the space character.

EM Implementation: the Baum–Welch algorithm Consider the hidden state H : $p_1 \dots p_n$ which we just initially guess - by choosing B - and the observation V : $c_1 \dots c_n$

- Each transition $p_{t-1} \rightarrow p_t$ happens with probability $P(p_t|p_{t-1})$. The c_i character is emitted with probability is $P(c_t|p_t)$.
- The probability of the visible sequence V given a particular hidden state sequence is the initial probability of starting with p_1 , $\Pi(p_1)$ and then emitting c_1 : $P(c_1|p_1)$, times the transition probabilities and the emission probabilities at each timestep t :

$$P(c_{1..N}, p_{1..N}|\Theta) = \Pi(p_1)P(c_1|p_1) \prod_t P(p_t|p_{t-1})P(c_t|p_t).$$
Notice that this can be computed recursively.
- The probability of the visible sequence V is the marginalization of the previous computation for every possible hidden sequence $H = p_{1..N}$: $P(V|\Theta) = \sum_H \Pi(p_1) \prod_i P(p_t|p_{t-1})P(c_t|p_t)$. This is a computational challenge.

- Baum-Welch eases the computation by splitting it in two parts for each timestep, based on *what has already been observed* (forward pass) and *what it will be observed* (backward pass).
- The **forward pass** of the Forward-Backward algorithm reduces the complexity by computing $P(H_t = p_t, V_{1\dots t} | \Theta)$ recursively and then marginalizing the result. At each step t for each hidden state i and visible output $V_t = c_t$ the value *alpha* is defined as follows:

$$\alpha_t[i] = B[V_t] \sum_j \alpha_{t-1}(j) A[ij]$$
- The **backward pass** will complete this partial result: where the forward pass computes the probability up to step t , the backward pass does the same for the subsequent timesteps, going backwards from step t to step $t+1$ and computing $P(V_{(t+1)\dots N} | p_t, \Theta)$. This pass can be again computed recursively in a similar way. For this computation the value *beta* is introduced:

$$\beta_t[i] = \sum_j \beta_{t+1}[j] B[j, V_{t+1}] A[ij]$$
- The probability of moving from state p_t to p_{t+1} is proportional to $\alpha_t[i] A[ij] B[j, V_{t+1}] \beta_{t+1}(j)$. Splitting this computation with alpha and beta is the core optimization that avoids the algorithm being exponential.
- The conditional probability of the hidden states H is: $P(H|V, \Theta) = \frac{P(H, V|\Theta)}{P(V|\Theta)}$. The partial results of the forward and backward passes can be later reused to calculate this probability.
- Finally, Π , A and B can be updated by marginalizing $P(H|V, \Theta)$ and normalizing.

Key implementation details (see [6]): - Transition matrix A entries stay non-zero to keep the Markov chain ergodic. - Forward/backward passes are scaled each step to prevent underflow on long ciphertexts. - First-order LM keeps sparsity manageable; higher orders need smoothing and more data.

EM implementation: a concrete example So far I have summarized the algorithm and its Forward-Backward passes. Let's use it to automatically decrypt some text.

EM usage sketch (see Appendix A for full code)

1. Build bigram LM (A, B_{id}, Π)[7]: `A, B_id, PI = build_bigram_model(corpus_text, max_iter=50, limit=200k).`
2. Initialize emission matrix B randomly, apply `constrain_space_mapping(B)`.
3. Run `BaumWelch.process(A, B, PI, V, max_iter=200, recompute_A=False, recompute_PI=False, ext_fun=constrain_space_mapping)`.
4. Decode with `decode_with_emissions(B_hat, ciphertext)` or feed emissions to Viterbi[8] for MAP decoding.

Log-likelihoods from a representative run are in the convergence table below.

EM convergence snapshot

| Iteration | $\log P(V \theta)$ | Note |
|-----------|--------------------|--|
| 0 | -374.15 | random init |
| 2 | -297.81 | structure starts to appear |
| 16 | -242.69 | readable plaintext emerges |
| 950 | -237.72 | plateau; rerun with new seeds to escape local optima |

Scaling prevents underflow, and zero-free A keeps the chain ergodic. Repeat with a few random B seeds and keep the run with the best log-likelihood.

EM decoded sample

- Iter 0: BZN BQWNXLN JQBZ ... (random initialization, gibberish)
- Iter 16: THE TOXUPLE WITH HAZIND AN OQEN MING ... (largely readable, a few bad mappings)
- Iter 950: THE TROUPLE WITH HAZING AN OQEN MIND ... (plateau; minor letter swaps remain)

These snapshots show how likelihood climbs and the mapping improves before stalling. Rerun with new seeds to escape plateaus.

EM usage example: computing statistics for the English bigrams LM Use `build_bigram_model` to estimate the bigram transition matrix A and initial state distribution Π from a cleaned English corpus (e.g., Brown). The helper trims the text to a configurable limit to keep runtimes reasonable.

EM usage: decrypt a text With A and Π fixed, initialize an emission matrix B , apply `constrain_space_mapping`, and run `BaumWelch.process` with `recompute_A=False` and `recompute_PI=False`. Use `decode_with_emissions` for a quick readout or plug the emissions into Viterbi[8] for a MAP sequence.

EM: Review of the Baum-Welch approach With the Baum-Welch algorithm I can automatically decrypt a text in seconds.

The algorithm is very flexible as any of its parameters A, B, Θ can be optimized. I used the algorithm to prepare the English bigram LM and used it again to decrypt some text.

let's make some considerations:

- the entropy of an English text tells us how much information is given us by a text. This can be used to compute the *unicity distance*: (Claude Shannon, “Communication Theory of Secrecy Systems”, 1949) [9]: the minimum length of text which gives us enough information to decrypt, which is defined as the entropy of the keyspace divided by the per-character redundancy in bits (ca. 3.2 for English text):

$$U = \frac{H(K)}{D} = \frac{\log_2(26!)}{3.2} \approx 28$$

For the Baum-Welch approach to work, much longer text is needed.

- this technique only exploits the bigram LM, or first-order HMM, which gives to the statistics a fair restricted horizon visibility.
- it's possible to improve the results by modifying the algorithm and make use of higher-order HMMs. Various solutions can be further used to improve speed and memory usage, but the algorithm will be slower and memory intensive when using higher-order HMMs.
- the substitution cipher realizes a bijective mapping, but I could not directly enforce nor exploit this knowledge in the algorithm.
- it's possible to give a hint to the process: like for the space-to-space mapping, the A matrix (randomly filled in the example) can hold higher probabilities for known or expected mappings. However, there is no immediate way to tell the process that there are high expectations for specific words of part of sentences, like greetings or names.

Break the cipher with Metropolis-Hastings

Let's try out other solutions for the problem.

Is there a way to enforce the bijective property of the cipher while solving the problem?

I briefly describe the Monte Carlo technique and the Metropolis-Hastings algorithm before using it with a new model.

Breaking the cipher with Expectation-Maximization

We cast the substitution cipher as a Hidden Markov Model: hidden plaintext states emit observed ciphertext symbols, with first-order transitions drawn from an English bigram LM. EM (Baum-Welch) tweaks the transition/emission probabilities to maximize the likelihood of the observed ciphertext, using scaled forward/backward passes to avoid underflow and non-zero transitions to keep the chain ergodic.

Monte Carlo

When we have a distribution which we can estimate, or from which can be sampled, if we want to find related quantities which for some reasons are intractable or not easy to compute, we can spare the calculations and just approximate these quantities by random sampling.

As a typical example, we could estimate the value of π by just drawing a circle of radius 1 inscribed in a 2x2 square and then repeatedly generate a pseudo-random point within the square area (coordinates $x, y \in [1, 1]$) and measure the proportion of times the point falls within the circle area.

Knowing the relation between the area of a square and the area of the inscribed circle:

$$\text{square area} = 2r * 2r = 4$$

$$\text{circle area} = \pi * r^2 = \pi$$

$$\frac{\text{circle area}}{\text{square area}} = \pi/4$$

We reach an estimation of the expected value of $\pi/4$ and, most importantly, for the law of large numbers this value will converge to the expected value $\pi/4$ for large n.

```
1 limit=1000000 # the more the better
2 print(sum([random.uniform(-1,1)**2 + random.uniform(-1,1)**2<=1 for i in range(limit)
    ])*4.0/limit)
3 3.14188
```

Under the “Monte Carlo” methods fall different techniques, which stochastically tweak the parameters of the model to explore unknown values and evaluate its outputs, such that the model can be improved. In a HMM this finetuning is called “random walk”.

This walk is mostly done by generating random samples and discarding not plausible ones (rejection sampling) or by generating samples from some weighted distribution (importance sampling).

For example, in the case of a substitution cipher, consider a random walk where at each step two letters of the cipher key are swapped and the resulting plaintext is sampled to evaluate its plausibility.

Metropolis-Hastings, brief introduction

Metropolis-Hastings is a Monte Carlo rejection sampling algorithm that works well for the substitution cipher case.

- consider an unknown distribution $P(x)$ which we however can estimate using only the observations.
- we have a function $P_l(x)$ which is just proportional to the density of $P(x)$.
- we make a random walk to progressively improve the unknown function by randomly make a change on the parameters of the model and then sample from the results.
- the new value will be accepted or refused using the “acceptance function” which compares $P_l(x)$ with its previous value to decide.

In the longer term (and under some basic conditions) this walk is guaranteed to optimize the system and return samples which follow the unknown distribution $P(x)$.

More about convergency of Markov models

A (irreducible, aperiodic) discrete Markov model with states distribution $\Pi(x)$ and transition distribution $P(y|x)$ will converge when:

$$\sum_x \Pi^*(x)P(y|x) = \Pi^*(y)$$

That is, the probability of being in y is equal to the probability of getting into y from any state x. Π^* is the stationary distribution here (it may not exist or may not be unique).

Starting with the initial state distribution Π and transitioning n times using the transition matrix A , the model will enter a new state: $\Pi * A^n$.

For high values of n , Π becomes insignificant and the result will converge:

$$\Pi^* = \lim_{n \rightarrow \infty} A^n \Rightarrow \Pi^* * A = \Pi^*$$

That is, if the probability Π has not changed after the transformation, the chain has converged (the random walk has ended, the system has not changed state).

A theorem says that if $A_{xy} > 0 \quad \forall x, y \rightarrow \exists$ unique Π^* (=if the transition matrix contains no zeroes, a unique stationary distribution exists).

The stationary distribution Π^* is equivalent to the unknown $P(x)$ and it's the desired result of the optimization process.

A useful way to see that the stationary distribution exists is when the system is time-reversible, meaning (*detailed balance equation*):

$$\Pi(x)P(y|x) = \Pi(y)P(x|y) \tag{3}$$

In fact, summing for x at both sides:

$$\sum_x \Pi(x)P(y|x) = \sum_x \Pi(y)P(x|y) = \Pi(y) \sum_x P(x|y) = \Pi(y)$$

Some clues behind the inner workings of Metropolis-Hastings for the substitution cipher

So far, the HMM has represented the probability of each individual letter in plain text to be transformed and observed. The transformation was not bijective and the first-order LM was a limitation.

Let's rethink the model in a way that enforces the substitution mapping and simplifies the use of higher-order LM.

The technique discussed here comes from the very interesting "The Markov Chain Monte Carlo Revolution" (Diaconis, Persi, 2009) [10] which provides deeper insights into MCMC and uses Metropolis-Hastings to, among other things, reveal the contents of an encrypted message from a prison inmate.

Consider (1) the substitution cipher mapping functions $S : f(P) = C$ and the inverse mapping $S^{-1} : f^{-1}(C) = P$ and its correspondence to the distribution functions $P(x|y)$ and $P(y|x)$ where X is the space of the possible symbols and $x, y \in X$

Then the stationary distribution Π would represent the probability that the mapping function is the correct one, as applying successive transformations will not change nor improve the model and the model will output a sample of this stationary distribution, which is the mapping of the cipher.

The Π distribution is initially unknown, but we can easily obtain a probability distribution which is proportional to this distribution.

In fact, using the statistics from the n-gram LM, to estimate the plausibility of a sentence as English text we can compute the joint probability P_{LM} of the n-grams in the text sequence $\{c_1, \dots, c_n\}$ transformed into candidate plaintext by the substitution S^{-1} :

$$(\text{Plausibility:}) \quad Pl = \prod_{i=1}^n P_{LM}(S^{-1}(c_{i+1}) | S^{-1}(c_i))$$

Metropolis-Hastings (symmetric proposal): the algorithm

The original Metropolis algorithm used the "symmetric proposal", which derives directly from the detailed balance equation.

It needs a symmetric function to choose a new candidate and another function to tune the results (some fitness function which in turn uses information from the LM) and it works as follows:

0. start with
 - some current state x (may be randomly generated)
 - an symmetrical distribution: $g(x|y)=g(y|x)$
 - a function $Pl(x)$ (the plausibility), proportional to the desired target $\Pi(x)$
1. from current x draw a new state y from the distribution $g(y|x)$
2. calculate the acceptance rate $\alpha(y, x) = \min(1, \frac{Pl(y)}{Pl(x)})$
3. Accept/Reject the candidate:
 - accept the candidate and move to y if $\alpha \geq \mu(0, 1)$ (not lower than a random uniform between 0 and 1)
 - refuse otherwise and stay in x
4. repeat from 1 until convergence

The acceptance ratio comes from (3):

$$\frac{P(y|x)}{P(x|y)} = \frac{\Pi(y)}{\Pi(x)}$$

using g to generate samples and alpha to accept or refuse the swap, rewrite the equation as:

$$\frac{g(x|y)\alpha(x, y)}{g(y|x)\alpha(y, x)} = \frac{\Pi(y)}{\Pi(x)}$$

because g is symmetric this simplifies to:

$$\frac{\alpha(x, y)}{\alpha(y, x)} = \frac{\Pi(y)}{\Pi(x)}$$

and because $Pl(x)$ is by requirement proportional to $\Pi(x)$ the condition can be fulfilled by choosing alpha to be:

$$\alpha(x, y) = \min\left(1, \frac{Pl(y)}{Pl(x)}\right) = \min\left(1, \frac{\Pi(y)}{\Pi(x)}\right)$$

```

1 cipher = normalize(ciphertext)
2 score = plausibility(cipher, lm)
3 for _ in range(N):
4     c1, c2 = pick_two_letters(cipher)
5     candidate = swap(cipher, c1, c2)Δ
6     = score(candidate) - score
7     accept with prob min(1, expΔ())
8     track best candidate for display

```

The proposal enforces a bijection by swapping two symbols; the plausibility function can mix multiple n-gram lengths.

Acceptance intuition: early iterations accept many swaps; as the score climbs, rejections dominate. Restart from a fresh seed if trapped in low-quality states.

Metropolis-Hastings implementation: a concrete example

With this algorithm, I am exploiting an LM with multiple n-grams. This is now easy due to the mild requirements for the Plausibility function. I don't even have to calculate a probability here. Consider that for each n-gram in the candidate text, the frequency of this n-gram in the plaintext will be higher if the n-gram has higher probabilities for that language. The sum of weighted n-gram frequencies (where weight is proportional to n-gram length as statistics on longer n-grams give us better information) is sufficient.

Actually, "Frequency" is a normalized quantity, I don't even need to normalize here as long as the normalization factor (the sum of n-gram counts) stays the same throughout the process, because at each timestep the normalization factor cancels out when the plausibilities Pl_t/Pl_{t+1} are compared.

Furthermore, at each step an existing mapping will be swapped, this maintains the bijection between Plaintext and Ciphertext alphabet.

The n-gram matrix representing the LM is sparse for larger n-gram lengths and it would be inefficient to simply store it in an array. In this python example I am using a *dictionary*.

Finally, the algorithm's convergence technique will allow it to jump from one local minimum to another. For this reason, there is no explicit way to determine whether it has completed. For my fairly simple purpose of breaking the cipher, I just set a maximum number of iterations; another possibility is to stop the algorithm if there is no improvement for a number of iterations.

Here is my full implementation of the Metropolis algorithm and a concrete example:

MH usage sketch (see Appendix B for full code)

- Build a multi-gram LM: `lm = count_ngrams(clean_corpus, min_len=2, max_len=5)`.
- Substitution cipher: `metropolis_substitution(ciphertext, lm, iterations=50k, seed=7)` prints improved candidates as it accepts better swaps.
- Patristocrat spacing: `metropolis_patristocrat(text_without_spaces, lm, iterations=5k, seed=11)` proposes insert/remove-space moves until the plausibility stabilizes.

Both samplers print intermediate best states so convergence can be inspected in the PDF output.

MH trace highlights

- Early iterations expose high-signal bigrams (e.g., TH, HE) and quickly improve readability.
- Acceptance rate is controlled by the plausibility spread; widening the n-gram range speeds convergence on short texts.
- Patristocrat spacing recovers steadily as high-frequency trigrams reappear.

MH trace sample

- It#1: K CXJO XI HBD BKIO XS ENJHB HEN XI HBD CLSB (initial guess)
- It#15: M CFJS FI TBR BMIS FO ENJTB TEN FI TBR CLOB (better bigrams emerging)
- It#219: IF YOU LOOK FOR PERFECTION YOU LL NEVER BE CONTENT (spacing recovered)

Selected iterations illustrate convergence without dumping full logs.

Patristocrats

To make the cipher more difficult to break, at the expense of possible ambiguities in the interpretation of the plaintext, all spaces may be removed from the text. The American Cryptogram Association calls this kind of substitution code "Patristocrat".

To solve a Patristocrat, we need an LM calculated from a corpus *without* spaces. The same algorithms as before can be used, but longer ciphertext is needed to provide sufficient statistical input.

The higher-order HMM implementations, such as the one used for Metropolis-Hastings, give here much better results than the first-order implementations, as expected.

After the plaintext has been restored, it is also possible to restore the spaces - this time using an LM from a corpus *with* spaces. There are of course several ways to achieve the goal, such as dynamic programming, but - with only minor changes from the previous implementation - here I use Metropolis for this:

For patristocrats (no spaces), build the language model without spaces and run `metropolis_patristocrat`. Spacing recovery leans on high-frequency trigrams and common prefixes/suffixes; ambiguity increases on very short texts.

Review of the Metropolis approach

Given enough time, very short text can be successfully decrypted with Metropolis.

The algorithm is quite flexible and allows the use of LM from higher order Markov models with ease.

The convergence rate is determined by the function that selects a new candidate and acceptance function (which determines the rejection rate), modified versions of these functions form the basis for alternate versions of this algorithm.

Limitations & future work

- Short ciphertexts remain ambiguous; try multiple seeds or add crib hints.
- Higher-order LMs could improve plausibility but need smoothing and larger corpora.
- Non-English alphabets/spaces require retraining LMs and adjusting symbol maps.
- EM can stall in local optima; multiple restarts or tempered EM could help.

Appendices

Full Python listings referenced in the text. Place these in `src/crypto/` or on your import path when re-running the experiments.

Appendix A: EM helper code

```
1 """
2 Baum-Welch routines tailored to simple substitution ciphers.
3 """
4
5 from __future__ import annotations
6
7 from typing import Callable, Tuple
8
9 import numpy as np
10
11 from .utils import (
12     ALPHABET,
13     SPACE_IDX,
14     char_to_index,
15     encode_text,
16     normalize_text,
17     seed_everything,
18 )
19
20 EmissionHook = Callable[[np.ndarray], np.ndarray]
21
22
23 class BaumWelch:
24     """Minimal Baum-Welch implementation with scaling to avoid underflow."""
25
26     def __init__(self, num_hidden: int = len(ALPHABET)):
27         self.num_hidden = num_hidden
28         self.num_visible = num_hidden
29         self.T = 0
30
31     @staticmethod
32     def normalize(matrix: np.ndarray) -> Tuple[np.ndarray, np.ndarray]:
33         totals = np.sum(matrix, axis=min(1, len(matrix.shape) - 1), keepdims=True)
```

```

34     totals[totals == 0] = 1.0
35     return matrix / totals
36
37 def alpha(self, A: np.ndarray, B: np.ndarray, PI: np.ndarray, V: np.ndarray):
38     alphas = np.zeros((self.T, self.num_hidden))
39     scale_facts = np.zeros(self.T)
40     alphas[0], scale_facts[0] = self.normalize(PI * B[:, V[0]])
41     for t in range(1, self.T):
42         alphas[t], scale_facts[t] = self.normalize(
43             B[:, V[t]] * np.sum(alphas[t - 1] * A.T, axis=1)
44         )
45     return alphas, scale_facts
46
47 def beta(self, A: np.ndarray, B: np.ndarray, V: np.ndarray, scale_facts: np.ndarray):
48     betas = np.zeros((self.T, self.num_hidden))
49     betas[self.T - 1] = np.array([1.0 / scale_facts[self.T - 1]])
50     for t in range(self.T - 2, -1, -1):
51         betas[t] = (
52             np.sum(betas[t + 1] * B[:, V[t + 1]] * A, axis=1) / scale_facts[t]
53         )
54     return betas
55
56 def process(
57     self,
58     A: np.ndarray,
59     B: np.ndarray,
60     PI: np.ndarray,
61     V: np.ndarray,
62     max_iter: int,
63     recompute_A: bool = True,
64     recompute_B: bool = True,
65     recompute_PI: bool = True,
66     ext_fun: EmissionHook | None = None,
67     verbose_fun: Callable[[np.ndarray, np.ndarray, int, float], None] | None = None,
68 ):
69     self.T = V.size
70     log_scale_facts = None
71     for it in range(max_iter):
72         alphas, scale_facts = self.alpha(A, B, PI, V)
73         new_log_scale = float(np.sum(np.log(scale_facts)))
74         if log_scale_facts is not None and new_log_scale <= log_scale_facts:
75             if verbose_fun:
76                 print("Algorithm has converged")
77                 break
78         log_scale_facts = new_log_scale
79         betas = self.beta(A, B, V, scale_facts)
80         gammas = np.zeros((self.T, self.num_hidden, self.num_hidden))
81         for t in range(self.T - 1):
82             gammas[t] = (A.T * alphas[t, :]).T * (betas[t + 1] * B[:, V[t + 1]])
83             gammas_i = np.sum(gammas, axis=2)
84             gammas_i[self.T - 1] = alphas[self.T - 1, :]
85
86         if recompute_A:
87             gammas_t = np.sum(gammas, axis=0)

```

```

88         totals = np.sum(gammas_t, axis=1)[:, None]
89         totals[totals == 0] = 1.0
90         A = (gammas_t.T / totals).T
91
92     if recompute_B:
93         for v in range(self.num_visible):
94             mask = V == v
95             if np.any(mask):
96                 B[:, v] = np.sum(gammas_i[mask], axis=0)
97         totals = np.sum(gammas_i, axis=0)
98         totals[totals == 0] = 1.0
99         B = (B.T / totals).T
100
101    if recompute_PI:
102        PI = gammas_i[0]
103
104    if ext_fun is not None:
105        B = ext_fun(B)
106
107    if verbose_fun:
108        verbose_fun(V, B, it, log_scale_facts)
109
110    return A, B, PI, log_scale_facts
111
112
113 def constrain_space_mapping(B: np.ndarray) -> np.ndarray:
114     """
115     Force a 1-1 mapping for space:space in the emission matrix.
116
117     Ensures space can only map to space, keeping the substitution mapping bijective for that
118     symbol.
119     """
120     B = B.copy()
121     B[SPACE_IDX, :] = 0.0
122     B[:, SPACE_IDX] = 0.0
123     B[SPACE_IDX, SPACE_IDX] = 1.0
124     row_sums = B.sum(axis=1, keepdims=True)
125     row_sums[row_sums == 0] = 1.0
126     return B / row_sums
127
128 def build_bigram_model(
129     corpus_text: str,
130     *,
131     max_iter: int = 150,
132     limit: int = 1_000_000,
133     seed: int = 13,
134 ):
135     """
136     Train a bigram language model from plain text using Baum-Welch.
137
138     Returns the transition matrix A, an identity emission matrix B, and the
139     initial state distribution PI. The text is upper-cased, cleaned, and
140     truncated to `limit` characters to keep runtimes reasonable.

```

```

141     """
142     seed_everything(seed)
143     normalized = normalize_text(corpus_text)[:limit]
144     V = np.array(encode_text(normalized))
145
146     num_states = len(ALPHABET)
147     PI = np.random.rand(num_states)
148     A = np.random.rand(num_states, num_states)
149     B = np.identity(num_states)
150
151     baum = BaumWelch(num_states)
152     A, _, PI, _ = baum.process(A, B, PI, V, max_iter=max_iter, recompute_B=False)
153     return A, B, PI
154
155
156 def decode_with_emissions(B: np.ndarray, observations: str) -> str:
157     """
158     Simple decoder: pick the highest emission probability per symbol.
159
160     For maximum a posteriori decoding, use the Viterbi algorithm instead.
161     """
162     V = [char_to_index(ch) for ch in observations]
163     decoded_indices = [int(np.argmax(B[:, v])) for v in V]
164     return ''.join(ALPHABET[idx] for idx in decoded_indices)

```

Appendix B: Metropolis-Hastings helper code

```

1 """
2 Metropolis-Hastings helpers for substitution and patristocrat variants.
3 """
4
5 from __future__ import annotations
6
7 import math
8 import random
9 import string
10 from collections import Counter
11 from typing import Dict, Iterable
12
13 from .utils import normalize_text
14
15
16 def count_ngrams(text: str, min_len: int = 2, max_len: int = 9) -> Dict[int, Dict[str, float]]:
17     """Count n-grams of varying length and return log counts for stability."""
18     return {
19         n: {k: math.log(v) for k, v in Counter(text[idx : idx + n] for idx in range(len(text) - n + 1)).items()}
20         for n in range(min_len, max_len + 1)
21     }
22
23
24 def plausibility_score(text: str, lm: Dict[int, Dict[str, float]]) -> float:
25     """Score text against an LM, weighted by n-gram length."""

```

```

26     counts = count_ngrams(text, min(lm), max(lm))
27     return sum(
28         lm[n].get(k, 0.0) * n
29         for n in counts
30         for k, _ in counts[n].items()
31         if k in lm[n]
32     )
33
34
35 def metropolis_substitution(
36     ciphertext: str,
37     lm: Dict[int, Dict[str, float]],
38     *,
39     iterations: int = 1_000_000,
40     seed: int = 13,
41 ) -> str:
42     """
43     Symmetric-proposal Metropolis-Hastings for simple substitution ciphers.
44
45     Fixed symbols (anything outside A-Z) are kept in place. The function yields
46     intermediate improvements to make convergence observable in static logs.
47     """
48     random.seed(seed)
49     ciphertext = normalize_text(ciphertext)
50     best = current_score = plausibility_score(ciphertext, lm)
51
52     fixed = set(ciphertext) - set(string.ascii_uppercase)
53     used_symbols = list(set(ciphertext) - fixed)
54     all_symbols = set(string.ascii_uppercase)
55     accepted = ciphertext
56     for i in range(iterations):
57         c1 = random.choice(used_symbols)
58         c2 = random.choice([c for c in all_symbols if c != c1])
59         candidate = accepted.translate(str.maketrans({c1: c2, c2: c1}))
60         candidate_score = plausibility_score(candidate, lm)
61         mu = math.log(random.uniform(0, 1))
62         if mu < max(min(candidate_score - current_score, 0), math.log(1e-3)):
63             accepted = candidate
64             current_score = candidate_score
65             used_symbols = list(set(accepted) - fixed)
66             if candidate_score > best:
67                 best = candidate_score
68                 print(f"It#{i}: {accepted}")
69     return accepted
70
71
72 def metropolis_patristocrat(
73     ciphertext: str,
74     lm: Dict[int, Dict[str, float]],
75     *,
76     iterations: int = 5000,
77     seed: int = 21,
78 ) -> str:
79     """

```

```

80     Restore spaces in patristocrat-style ciphertexts using a MH sampler.
81
82     The proposal either removes a space or inserts one adjacent to letters. A
83     length-normalized plausibility score keeps proposals comparable.
84     """
85     random.seed(seed)
86     ciphertext = normalize_text(ciphertext).replace(" ", " ")
87     best = current = ciphertext
88     best_score = current_score = plausibility_score(ciphertext, lm) / len(ciphertext)
89
90     for i in range(iterations):
91         c = random.choice(range(1, len(current) - 1))
92         candidate = list(current)
93         if current[c].isspace():
94             candidate.pop(c)
95         elif current[c - 1] != " " and current[c + 1] != " ":
96             candidate.insert(c, " ")
97         candidate_text = "".join(candidate)
98         cand_score = plausibility_score(candidate_text, lm) / len(candidate_text)
99         mu = math.log(random.uniform(0, 1))
100        if mu < min(cand_score - current_score, 0):
101            current = candidate_text
102            current_score = cand_score
103            if cand_score > best_score:
104                best = current
105                best_score = cand_score
106                print(best)
107    return best

```

Appendix C: Utility functions

```

1 import random
2 import re
3 from typing import Iterable
4
5 ALPHABET = "ABCDEFGHIJKLMNOPQRSTUVWXYZ "
6 SPACE_IDX = ALPHABET.index(" ")
7
8
9 def seed_everything(seed: int) -> None:
10     """Seed Python's random generator and numpy if available."""
11     random.seed(seed)
12     try:
13         import numpy as np
14
15         np.random.seed(seed)
16     except Exception:
17         # numpy is optional
18         pass
19
20
21 def normalize_text(text: str) -> str:
22     """
23     Keep A-Z and spaces, collapse whitespace, and uppercase.

```

```

24
25     Ensures corpus and cipher text share the same character mapping for the LM.
26     """
27     cleaned = re.sub("[^A-Z ]", " ", text.upper())
28     return re.sub("\s+", " ", cleaned).strip()
29
30
31 def char_to_index(ch: str) -> int:
32     """Map a character in ALPHABET to its ordinal index."""
33     return ALPHABET.index(ch)
34
35
36 def index_to_char(idx: int) -> str:
37     """Map an ordinal index back to a character in ALPHABET."""
38     return ALPHABET[idx]
39
40
41 def encode_text(text: Iterable[str]) -> list[int]:
42     """Convert a string of characters into indices."""
43     return [char_to_index(ch) for ch in text]
44
45
46 def decode_indices(indices: Iterable[int]) -> str:
47     """Convert indices back to a string of characters."""
48     return "".join(index_to_char(i) for i in indices)

```

Appendix D: Package exports

```

1 """
2 Lightweight cryptanalytic helpers for EM- and MH-based substitution cipher exploration.
3 """
4
5 from .utils import (
6     ALPHABET,
7     SPACE_IDX,
8     char_to_index,
9     index_to_char,
10    encode_text,
11    decode_indices,
12    normalize_text,
13    seed_everything,
14 )
15 from .em import BaumWelch, build_bigram_model, constrain_space_mapping
16 from .mh import (
17     count_ngrams,
18     plausibility_score,
19     metropolis_substitution,
20     metropolis_patristocrat,
21 )
22
23 __all__ = [
24     "ALPHABET",
25     "SPACE_IDX",
26     "char_to_index",

```

```

27     "index_to_char",
28     "encode_text",
29     "decode_indices",
30     "normalize_text",
31     "seed_everything",
32     "BaumWelch",
33     "build_bigram_model",
34     "constrain_space_mapping",
35     "count_ngrams",
36     "plausibility_score",
37     "metropolis_substitution",
38     "metropolis_patristocrat",
39 ]

```

References

- [1] Substitution cipher. (2021, July 27). In Wikipedia. https://en.wikipedia.org/w/index.php?title=Substitution_cipher&oldid=1035776944
- [2] Dempster, A.P.; Laird, N.M.; Rubin, D.B. (1977). “Maximum Likelihood from Incomplete Data via the EM Algorithm”. *Journal of the Royal Statistical Society, Series B*. 39 (1): 1–38. JSTOR 2984875. MR 0501537.
- [3] Maya R. Gupta and Yihua Chen (2011), “Theory and Use of the EM Algorithm”, *Foundations and Trends® in Signal Processing*: Vol. 4: No. 3, pp 223-296. <https://dx.doi.org/10.1561/2000000034>
- [4] Knight, Kevin & Nair, Anish & Rathod, Nishit & Yamada, Kenji. (2006). Unsupervised Analysis for Decipherment Problems. <https://dx.doi.org/10.3115/1273073.1273138>.
- [5] Gagniuc, Paul A. (2017). *Markov Chains: From Theory to Implementation and Experimentation*. USA, NJ: John Wiley & Sons. pp. 1–256. ISBN 978-1-119-38755-8
- [6] L. R. Rabiner, “A tutorial on hidden Markov models and selected applications in speech recognition,” in Proceedings of the IEEE, vol. 77, no. 2, pp. 257-286, Feb. 1989. <https://dx.doi.org/10.1109/5.18626>.
- [7] Francis, W. Nelson & Henry Kucera. 1967. Computational Analysis of Present-Day American English. Providence, RI: Brown University Press.
- [8] A. Viterbi, “Error bounds for convolutional codes and an asymptotically optimum decoding algorithm,” in IEEE Transactions on Information Theory, vol. 13, no. 2, pp. 260-269, April 1967, doi: 10.1109/TIT.1967.1054010.
- [9] Shannon, Claude. “Communication Theory of Secrecy Systems”, *Bell System Technical Journal*, vol. 28(4), page 656–715, 1949. <https://doi.org/10.1002/j.1538-7305.1949.tb00928.x>
- [10] Diaconis, Persi. (2009). The Markov Chain Monte Carlo Revolution. *Bulletin of the American Mathematical Society*. 46. 179–205. <https://dx.doi.org/10.1090/S0273-0979-08-01238-X>.